

Application Layer

CS446 / CS646 - Networking
Instructor: Bo Sheng



1

Application layer: overview

- Principles of network applications
- Web and HTTP
- E-mail, SMTP, IMAP
- The Domain Name System DNS
- P2P applications
- video streaming and content distribution networks
- socket programming with UDP and TCP



2

Application layer: overview

Our goals:

- conceptual *and* implementation aspects of application-layer protocols
 - transport-layer service models
 - client-server paradigm
 - peer-to-peer paradigm
- learn about protocols by examining popular application-layer protocols and infrastructure
 - HTTP
 - SMTP, IMAP
 - DNS
 - video streaming systems, CDNs
- programming network applications
 - socket API



3

Some network apps

- social networking
- Web
- text messaging
- e-mail
- multi-user network games
- streaming stored video (YouTube, Hulu, Netflix)
- P2P file sharing
- voice over IP (e.g., Skype)
- real-time video conferencing (e.g., Zoom)
- Internet search
- remote login
- ...



4


Creating a network app

write programs that:

- run on (different) end systems
- communicate over network
- e.g., web server software communicates with browser software

no need to write software for network-core devices

- network-core devices do not run user applications
- applications on end systems allows for rapid app development, propagation



5

Client-server paradigm


server:

- always-on host
- permanent IP address
- often in data centers, for scaling

clients:

- contact, communicate with server
- may be intermittently connected
- may have dynamic IP addresses
- do *not* communicate directly with each other


- examples: HTTP, IMAP, FTP



6

Peer-peer architecture

- no always-on server
- arbitrary end systems directly communicate
- peers request service from other peers, provide service in return to other peers
 - *self scalability* – new peers bring new service capacity, as well as new service demands
- peers are intermittently connected and change IP addresses
 - complex management
- example: P2P file sharing



Processes communicating

process: program running within a host

- within same host, two processes communicate using **inter-process communication** (defined by OS)
- processes in different hosts communicate by exchanging **messages**

- clients, servers

- client process:** process that initiates communication
- server process:** process that waits to be contacted

- note: applications with P2P architectures have client processes & server processes



Sockets

- process sends/receives messages to/from its **socket**
- socket analogous to door
 - sending process shoves message out the door
 - sending process relies on transport infrastructure on the other side of the door to deliver the message to the socket at the receiving process
 - two sockets involved: one on each side



Addressing processes

- to receive messages, process must have **identifier**
- host device has unique 32-bit IP address
- Q: does IP address of host on which process runs suffice for identifying the process?
- A: no, *many* processes can be running on same host
- **identifier** includes both **IP address** and **port numbers** associated with process on host.
- example port numbers:
 - HTTP server: 80
 - mail server: 25
- to send HTTP message to `gaia.cs.umass.edu` web server:
 - IP address: 128.119.245.12
 - port number: 80
- more shortly...



An application-layer protocol defines:

- **types of messages exchanged,**
 - e.g., request, response
- **message syntax:**
 - what fields in messages & how fields are delineated
- **message semantics**
 - meaning of information in fields
- **rules** for when and how processes send & respond to messages

open protocols:

- defined in RFCs, everyone has access to protocol definition
- allows for interoperability
- e.g., HTTP, SMTP

proprietary protocols:

- e.g., Skype, Zoom



What transport service does an app need?

data integrity

- some apps (e.g., file transfer, web transactions) require 100% reliable data transfer
- other apps (e.g., audio) can tolerate some loss

timing

- some apps (e.g., Internet telephony, interactive games) require low delay to be “effective”

throughput

- some apps (e.g., multimedia) require minimum amount of throughput to be "effective"
- other apps ("elastic apps") make use of whatever throughput they get

security

- encryption, data integrity, ...



Transport service requirements: common apps

application	data loss	throughput	time sensitive?
file transfer/download	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic	no
real-time audio/video	loss-tolerant	audio: 5Kbps-1Mbps video: 10Kbps-5Mbps	yes, 10's msec
streaming audio/video	loss-tolerant	same as above	yes, few secs
interactive games	loss-tolerant	Kbps+	yes, 10's msec
text messaging	no loss	elastic	yes and no



13

Internet transport protocols services

TCP service:

- **reliable transport** between sending and receiving process
- **flow control**: sender won't overwhelm receiver
- **congestion control**: throttle sender when network overloaded
- **connection-oriented**: setup required between client and server processes
- **does not provide**: timing, minimum throughput guarantee, security

UDP service:

- **unreliable data transfer** between sending and receiving process
- **does not provide**: reliability, flow control, congestion control, timing, throughput guarantee, security, or connection setup.

Q: why bother? Why is there a UDP?



14

Internet applications, and transport protocols

application	application layer protocol	transport protocol
file transfer/download	FTP [RFC 959]	TCP
e-mail	SMTP [RFC 5321]	TCP
Web documents	HTTP 1.1 [RFC 7320]	TCP
Internet telephony	SIP [RFC 3261], RTP [RFC 3550], or proprietary	TCP or UDP
streaming audio/video	HTTP [RFC 7320], DASH	TCP
interactive games	WOW, FPS (proprietary)	UDP or TCP



15

Securing TCP

Vanilla TCP & UDP sockets:

- no encryption
- cleartext passwords sent into socket traverse Internet in cleartext (!)

Transport Layer Security (TLS)

- provides encrypted TCP connections
- data integrity
- end-point authentication

TSL implemented in application layer

- apps use TSL libraries, that use TCP in turn
- cleartext sent into "socket" traverse Internet *encrypted*
- more: Chapter 8



16

Application layer: overview

- Principles of network applications
- **Web and HTTP**
- E-mail, SMTP, IMAP
- The Domain Name System DNS

- P2P applications
- video streaming and content distribution networks
- socket programming with UDP and TCP



17

Web and HTTP

First, a quick review...

- web page consists of **objects**, each of which can be stored on different Web servers
- object can be HTML file, JPEG image, Java applet, audio file,...
- web page consists of **base HTML-file** which includes **several referenced objects**, each addressable by a **URL**, e.g.,

www.someschool.edu/someDept/pic.gif
host name
path name



18

HTTP overview

HTTP: hypertext transfer protocol

- Web's application-layer protocol

client/server model:

- client:** browser that requests, receives, (using HTTP protocol) and "displays" Web objects
- server:** Web server sends (using HTTP protocol) objects in response to requests



19

HTTP overview (continued)

HTTP uses TCP:

- client initiates TCP connection (creates socket) to server, port 80
- server accepts TCP connection from client
- HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- TCP connection closed

HTTP is "stateless"

- server maintains *no* information about past client requests

aside protocols that maintain "state" are complex!

- past history (state) must be maintained
- if server/client crashes, their views of "state" may be inconsistent, must be reconciled

20

HTTP connections: two types

Non-persistent HTTP

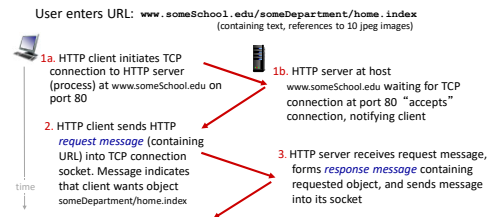
- TCP connection opened
 - at most one object sent over TCP connection
 - TCP connection closed
- downloading multiple objects required multiple connections

Persistent HTTP

- TCP connection opened to a server
- multiple objects can be sent over *single* TCP connection between client, and that server
- TCP connection closed

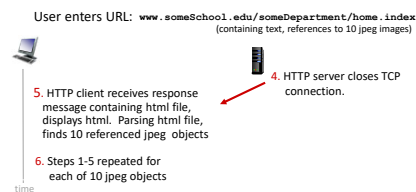
21

Non-persistent HTTP: example



22

Non-persistent HTTP: example (cont.)




23

Non-persistent HTTP: response time

RTT (definition): time for a small packet to travel from client to server and back

HTTP response time (per object):

- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
- object/file transmission time



Non-persistent HTTP response time = $2RTT + \text{file transmission time}$

24

Persistent HTTP (HTTP 1.1)

Non-persistent HTTP issues:

- requires 2 RTTs per object
- OS overhead for *each* TCP connection
- browsers often open multiple parallel TCP connections to fetch referenced objects in parallel

Persistent HTTP (HTTP1.1):

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects (cutting response time in half)

25

HTTP request message

- two types of HTTP messages: *request, response*

HTTP request message:

- ASCII (human-readable format)


request line (GET, POST, HEAD commands)

carriage return character
line-feed character

carriage return, line feed
at start of line indicates
end of header lines

26

HTTP request message: general format



27

Other HTTP request messages

POST method:

- web page often includes form input
- user input sent from client to server in entity body of HTTP POST request message

HEAD method:

- requests headers (only) that would be returned if specified URL were requested with an HTTP GET method.

PUT method:

- uploads new file (object) to server
- completely replaces file that exists at specified URL with content in entity body of POST HTTP request message

GET method (for sending data to server):

- include user data in URL field of HTTP GET request message (following a "?"):

www.scmesite.com/animalsearch?monkey&banana

28

HTTP response message

status line (protocol status code status phrase) → HTTP/1.1 200 OK

29

HTTP response status codes

- status code appears in 1st line in server-to-client response message.
- some sample codes:

200 OK

- request succeeded, requested object later in this message

301 Moved Permanently

- requested object moved, new location specified later in this message (in Location: field)

400 Bad Request

- request msg not understood by server

404 Not Found

- requested document not found on this server

505 HTTP Version Not Supported

30

Trying out HTTP (client side) for yourself

1. netcat to your favorite Web server:
 - `% nc -c -v gaia.cs.umass.edu 80`
 - opens TCP connection to port 80 (default HTTP server port) at `gaia.cs.umass.edu`.
 - anything typed in will be sent to port 80 at `gaia.cs.umass.edu`
2. type in a GET HTTP request:


```
GET /kucrose_ross/interactive/index.php HTTP/1.1
Host: gaia.cs.umass.edu
```

 - by typing this in (hit carriage return twice), you send this minimal (but complete) GET request to HTTP server
3. look at response message sent by HTTP server!




31

Maintaining user/server state: cookies

Recall: HTTP GET/response interaction is *stateless*

- no notion of multi-step exchanges of HTTP messages to complete a Web "transaction"
 - no need for client/server to track "state" of multi-step exchange
- all HTTP requests are independent of each other
 - no need for client/server to "recover" from a partially-completed-but-never-completely-completed transaction



32

Maintaining user/server state: cookies

Web sites and client browser use *cookies* to maintain some state between transactions

four components:

- 1) cookie header line of HTTP *response* message
- 2) cookie header line in next HTTP *request* message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site


Example:

- Susan uses browser on laptop, visits specific e-commerce site for first time
 - when initial HTTP requests arrives at site, site creates:
 - unique ID (aka "cookie")
 - entry in backend database for ID
- subsequent HTTP requests from Susan to this site will contain cookie ID value, allowing site to "identify" Susan



33

Maintaining user/server state: cookies



34

HTTP cookies: comments

What cookies can be used for:

- authorization
- shopping carts
- recommendations
- user session state (Web e-mail)

Challenge: How to keep state?

- at *protocol endpoints*: maintain state at sender/receiver over multiple transactions
- in *messages*: cookies in HTTP messages carry state

aside
cookies and privacy:

- cookies permit sites to learn a lot about you on their site.
- third party persistent cookies (tracking cookies) allow common identity (cookie value) to be tracked across multiple web sites




35

Web caches

Goal: satisfy client requests without involving origin server

- user configures browser to point to a (local) *Web cache*
- browser sends all HTTP requests to cache
 - if object in cache: cache returns object to client
 - else cache requests object from origin server, caches received object, then returns object to client



36

Web caches (aka proxy servers)

- Web cache acts as both client and server
 - server for original requesting client
 - client to origin server
 - server tells cache about object's allowable caching in response header:


```
Cache-Control: max-age=<seconds>
```

```
Cache-Control: no-cache
```
- Why Web caching?**
- reduce response time for client request
 - cache is closer to client
 - reduce traffic on an institution's access link
 - Internet is dense with caches
 - enables "poor" content providers to more effectively deliver content

37

Caching example

- Scenario:**
- access link rate: 1.54 Mbps
 - RTT from institutional router to server: 2 sec
 - web object size: 100K bits
 - average request rate from browsers to origin servers: 15/sec
 - avg data rate to browsers: 1.50 Mbps
- Performance:**
- access link utilization = .97 *problem: large queuing delays at high utilization!*
 - LAN utilization: .0015
 - end-end delay = Internet delay + access link delay + LAN delay
 - = 2 sec + minutes + usecs
-

38

Option 1: buy a faster access link

- Scenario:**
- access link rate: 154 Mbps
 - RTT from institutional router to server: 2 sec
 - web object size: 100K bits
 - average request rate from browsers to origin servers: 15/sec
 - avg data rate to browsers: 1.50 Mbps
- Performance:**
- access link utilization = .0097
 - LAN utilization: .0015
 - end-end delay = Internet delay + access link delay + LAN delay
 - = 2 sec + usecs + usecs
- Cost:** faster access link (expensive!)
-

39

Option 2: install a web cache

- Scenario:**
- access link rate: 1.54 Mbps
 - RTT from institutional router to server: 2 sec
 - web object size: 100K bits
 - average request rate from browsers to origin servers: 15/sec
 - avg data rate to browsers: 1.50 Mbps
- Cost:** web cache (cheap!)
- Performance:**
- LAN utilization: ?
 - access link utilization = ? *How to compute link utilization, delay?*
 - average end-end delay = ?
-

40

Calculating access link utilization, end-end delay with cache:

- suppose cache hit rate is 0.4:
- 40% requests served by cache, with low (msec) delay
 - 60% requests satisfied at origin
 - rate to browsers over access link = $0.6 * 1.50 \text{ Mbps} = .9 \text{ Mbps}$
 - access link utilization = $0.9 / 1.54 = .58$ means low (msec) queuing delay at access link
 - average end-end delay:
 - = $0.6 * (\text{delay from origin servers})$
 - + $0.4 * (\text{delay when satisfied at cache})$
 - = $0.6 (2.01) + 0.4 (\sim \text{msecs}) \approx \sim 1.2 \text{ secs}$
- lower average end-end delay than with 154 Mbps link (and cheaper too!)*
-


41

Conditional GET

Goal: don't send object if cache has up-to-date cached version

- no object transmission delay (or use of network resources)

- client:** specify date of cached copy in HTTP request
 - If-modified-since: <date>
- server:** response contains no object if cached copy is up-to-date:
 - HTTP/1.0 304 Not Modified



42

HTTP/2

Key goal: decreased delay in multi-object HTTP requests

HTTP/1.1: introduced multiple, pipelined GETs over single TCP connection

- server responds *in-order* (FCFS: first-come-first-served scheduling) to GET requests
- with FCFS, small object may have to wait for transmission (**head-of-line (HOL) blocking**) behind large object(s)
- loss recovery (retransmitting lost TCP segments) stalls object transmission



43

HTTP/2

Key goal: decreased delay in multi-object HTTP requests

HTTP/2: [RFC 7540, 2015] increased flexibility at *server* in sending objects to client:


- methods, status codes, most header fields unchanged from HTTP 1.1
- transmission order of requested objects based on client-specified object priority (not necessarily FCFS)
- push* unrequested objects to client
- divide objects into frames, schedule frames to mitigate HOL blocking



44

HTTP/2: mitigating HOL blocking


HTTP 1.1: client requests 1 large object (e.g., video file) and 3 smaller objects



45

HTTP/2: mitigating HOL blocking

HTTP/2: objects divided into frames, frame transmission interleaved



46

HTTP/2 to HTTP/3

HTTP/2 over single TCP connection means:

- recovery from packet loss still stalls all object transmissions
 - as in HTTP 1.1, browsers have incentive to open multiple parallel TCP connections to reduce stalling, increase overall throughput
- no security over vanilla TCP connection
- HTTP/3:** adds security, per object error- and congestion-control (more pipelining) over UDP
 - more on HTTP/3 in transport layer



47

Outline

- principles of network applications
- Web and HTTP
- electronic mail**
- DNS
- P2P applications
- video streaming and content distribution networks



48

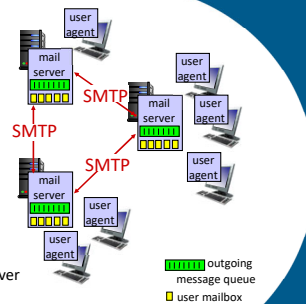
E-mail

Three major components:

- user agents
- mail servers
- simple mail transfer protocol: SMTP

User Agent

- a.k.a. "mail reader"
- composing, editing, reading mail messages
- e.g., Outlook, iPhone mail client
- outgoing, incoming messages stored on server



49

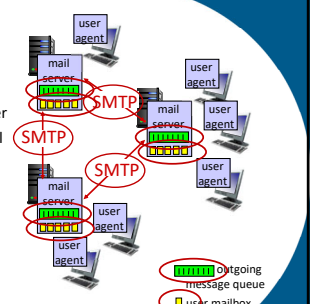
E-mail: mail servers

mail servers:

- mailbox contains incoming messages for user
- message queue of outgoing (to be sent) mail messages

SMTP protocol between mail servers to send email messages

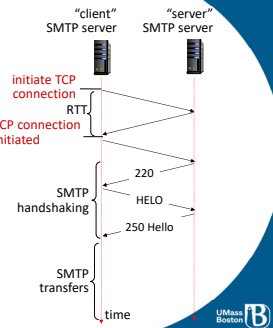
- client: sending mail server
- "server": receiving mail server



50

SMTP RFC (5321)


- uses TCP to reliably transfer email message from client (mail server initiating connection) to server, port 25
 - direct transfer: sending server (acting like client) to receiving server
- three phases of transfer
 - SMTP handshaking (greeting)
 - SMTP transfer of messages
 - SMTP closure
- command/response interaction (like HTTP)
 - commands: ASCII text
 - response: status code and phrase



51

Scenario: Alice sends e-mail to Bob

- Alice uses UA to compose e-mail message "to" bob@someschool.edu
- Alice's UA sends message to her mail server using SMTP; message placed in message queue
- client side of SMTP at mail server opens TCP connection with Bob's mail server
- SMTP client sends Alice's message over the TCP connection
- Bob's mail server places the message in Bob's mailbox
- Bob invokes his user agent to read message



52

Sample SMTP interaction

```
S: 220 hamburger.edu
C: HELO crepes.fr
S: 250 Hello crepes.fr, pleased to meet you
C: MAIL FROM: <alice@crepes.fr>
S: 250 alice@crepes.fr... Sender ok
C: RCPT TO: <bob@hamburger.edu>
S: 250 bob@hamburger.edu ... Recipient ok
C: DATA
S: 354 Enter mail, end with "." on a line by itself
C: Do you like ketchup?
C: How about pickles?
C: .
S: 250 Message accepted for delivery
C: QUIT
S: 221 hamburger.edu closing connection
```

53

Try SMTP interaction for yourself:

- telnet servername 25
- see 220 reply from server
- enter HELO, MAIL FROM, RCPT TO, DATA, QUIT commands

above lets you send email without using email client (reader)

54

Try SMTP interaction for yourself:

```
telnet mx1.cs.umb.edu 25
HELO cs.umb.edu
MAIL from: a@cs.umb.edu
RCPT to: b@cs.umb.edu
DATA
Subject: Test

This is a test
.
```

55

SMTP: observations

comparison with HTTP:

- HTTP: client pull
- SMTP: client push
- both have ASCII command/response interaction, status codes
- HTTP: each object encapsulated in its own response message
- SMTP: multiple objects sent in multipart message
- SMTP uses persistent connections
- SMTP requires message (header & body) to be in 7-bit ASCII
- SMTP server uses CRLF.CRLF to determine end of message


56

Mail message format

SMTP: protocol for exchanging e-mail messages, defined in RFC 5321 (like RFC 7231 defines HTTP)


RFC 2822 defines *syntax* for e-mail message itself (like HTML defines syntax for web documents)

- header lines, e.g.,
 - To:
 - From:
 - Subject:
 these lines, within the body of the email message area different from SMTP MAIL FROM:, RCPT TO: commands!
- Body: the "message", ASCII characters only



57

Retrieving email: mail access protocols



- **SMTP**: delivery/storage of e-mail messages to receiver's server
- mail access protocol: retrieval from server
 - **IMAP**: Internet Mail Access Protocol [RFC 3501]: messages stored on server, IMAP provides retrieval, deletion, folders of stored messages on server
 - **HTTP**: gmail, Hotmail, Yahoo!Mail, etc. provides web-based interface on top of SMTP (to send), IMAP (or POP) to retrieve e-mail messages

58

POP3 protocol

authorization phase

- client commands:
 - **user**: declare username
 - **pass**: password
- server responses:
 - +OK
 - -ERR

transaction phase, client:

- **list**: list message numbers
- **retr**: retrieve message by number
- **dele**: delete
- **quit**

```
S: +OK POP3 server ready
C: user bob
S: +OK
C: pass hungry
S: +OK user successfully logged on

C: list
S: 1 498
S: 2 912
S: .
C: retr 1
S: <message 1 contents>
S: .
C: dele 1
C: retr 2
S: <message 1 contents>
S: .
C: dele 2
C: quit
S: +OK POP3 server signing off
```

59

POP3 (more) and IMAP

more about POP3

- previous example uses POP3 "download and delete" mode
 - Bob cannot re-read e-mail if he changes client
- POP3 "download-and-keep": copies of messages on different clients
- POP3 is stateless across sessions

IMAP

- keeps all messages in one place: at server
- allows user to organize messages in folders
- keeps user state across sessions:
 - names of folders and mappings between message IDs and folder name

60

Application Layer: Overview

- Principles of network applications
- Web and HTTP
- E-mail, SMTP, IMAP
- The Domain Name System DNS
- P2P applications
- video streaming and content distribution networks
- socket programming with UDP and TCP



61

DNS: Domain Name System

people: many identifiers:

- SSN, name, passport #

Internet hosts, routers:

- IP address (32 bit) - used for addressing datagrams
- "name", e.g., cs.umass.edu - used by humans

Q: how to map between IP address and name, and vice versa ?

Domain Name System (DNS):

- distributed database** implemented in hierarchy of many **name servers**
- application-layer protocol:** hosts, DNS servers communicate to **resolve** names (address/name translation)
 - note:** core Internet function, **implemented as application-layer protocol**
 - complexity at network's "edge"

62

DNS: services, structure

DNS services:

- hostname-to-IP-address translation
- host aliasing
 - canonical, alias names
- mail server aliasing
- load distribution
 - replicated Web servers: many IP addresses correspond to one name

Q: Why not centralize DNS?

- single point of failure
- traffic volume
- distant centralized database
- maintenance

A: doesn't scale!

- Comcast DNS servers alone: 600B DNS queries/day
- Akamai DNS servers alone: 2.2T DNS queries/day

63

Thinking about the DNS

humongous distributed database:

- ~ billion records, each simple

handles many trillions of queries/day:

- many more reads than writes
- performance matters:** almost every Internet transaction interacts with DNS - msec count!

organizationally, physically decentralized:


- millions of different organizations responsible for their records

"bulletproof": reliability, security



64

DNS: a distributed, hierarchical database




Client wants IP address for www.amazon.com; 1st approximation:

- client queries root server to find .com DNS server
- client queries .com DNS server to get amazon.com DNS server
- client queries amazon.com DNS server to get IP address for www.amazon.com

65

DNS: root name servers

- official, contact-of-last-resort by name servers that can not resolve name



66

DNS: root name servers

- official, contact-of-last-resort by name servers that can not resolve name
- incredibly important** Internet function
 - Internet couldn't function without it!
 - DNSSEC – provides security (authentication, message integrity)
- ICANN (Internet Corporation for Assigned Names and Numbers) manages root DNS domain

13 logical root name “servers” worldwide each “server” replicated many times (~200 servers in US)

UMass Boston

67

Top-Level Domain, and authoritative servers

Top-Level Domain (TLD) servers:

- responsible for .com, .org, .net, .edu, .aero, .jobs, .museums, and all top-level country domains, e.g.: .cn, .uk, .fr, .ca, .jp
- Network Solutions: authoritative registry for .com, .net TLD
- Educause: .edu TLD

authoritative DNS servers:

- organization's own DNS server(s), providing authoritative hostname to IP mappings for organization's named hosts
- can be maintained by organization or service provider

UMass Boston

68

Local DNS name servers

- when host makes DNS query, it is sent to its *local* DNS server
 - Local DNS server returns reply, answering:
 - from its local cache of recent name-to-address translation pairs (possibly out of date!)
 - forwarding request into DNS hierarchy for resolution
 - each ISP has local DNS name server; to find yours:
 - MacOS: `% scutil --dns`
 - Windows: `>ipconfig /all`
- local DNS server doesn't strictly belong to hierarchy

UMass Boston

69

DNS name resolution: iterated query

Example: host at engineering.nyu.edu wants IP address for gaia.cs.umass.edu

Iterated query:

- contacted server replies with name of server to contact
- "I don't know this name, but ask this server"

UMass Boston

70

DNS name resolution: recursive query

Example: host at engineering.nyu.edu wants IP address for gaia.cs.umass.edu

Recursive query:

- puts burden of name resolution on contacted name server
- heavy load at upper levels of hierarchy?

UMass Boston

71

Caching DNS Information

- once (any) name server learns mapping, it *caches* mapping, and *immediately* returns a cached mapping in response to a query
 - caching improves response time
 - cache entries timeout (disappear) after some time (TTL)
 - TLD servers typically cached in local name servers
- cached entries may be *out-of-date*
 - if named host changes IP address, may not be known Internet-wide until all TTLs expire!
 - best-effort name-to-address translation!*

UMass Boston

72

Example

```
/etc/resolv.conf

dig google.com

dig cs.umb.edu      dig cs.umb.edu @8.8.8.8

dig cs.umb.edu MX / NS

dig cnnn.com

dig +norecurs google.com
```

73

DNS records

DNS: distributed database storing resource records (**RR**)

RR format: (name, value, type, ttl)

type=A

- name is hostname
- value is IP address

type=NS

- name is domain (e.g., foo.com)
- value is hostname of authoritative name server for this domain

type=CNAME

- name is alias name for some "canonical" (the real) name
- www.ibm.com is really servereast.backup2.ibm.com
- value is canonical name

type=MX

- value is name of mailserver associated with name

74

DNS protocol, messages

- query and reply messages, both with same message format

message header

- identification: 16 bit # for query, reply to query uses same #
- flags:
 - query or reply
 - recursion desired
 - recursion available
 - reply is authoritative

2 bytes		2 bytes	
identification	flags	# questions	# answer RRs
# authority RRs	# additional RRs	questions (variable # of questions)	answers (variable # of RRs)
authority (variable # of RRs)	additional info (variable # of RRs)		

75

DNS protocol, messages

name, type fields for a query

RRs in response to query

records for authoritative servers

additional "helpful" info that may be used

2 bytes		2 bytes	
identification	flags	# questions	# answer RRs
# authority RRs	# additional RRs	questions (variable # of questions)	answers (variable # of RRs)
authority (variable # of RRs)	additional info (variable # of RRs)		

76

Inserting records into DNS

- example: new startup "Network Utopia"
- register name networkutopia.com at **DNS registrar** (e.g., Network Solutions)
 - provide names, IP addresses of authoritative name server (primary and secondary)
 - registrar inserts two RRs into .com TLD server:
 - (networkutopia.com, dns1.networkutopia.com, NS)
 - (dns1.networkutopia.com, 212.212.212.1, A)
- create authoritative server type A record for www.networkutopia.com; type MX record for networkutopia.com

77

Attacking DNS

DDoS attacks

- bombard root servers with traffic
 - not successful to date
 - traffic filtering
 - local DNS servers cache IPs of TLD servers, allowing root server bypass
- bombard TLD servers
 - potentially more dangerous

redirect attacks

- man-in-middle
 - Intercept queries
- DNS poisoning
 - Send bogus replies to DNS server, which caches

exploit DNS for DDoS

- send queries with spoofed source address: target IP
- requires amplification

78

Outline

- principles of network applications
- Web and HTTP
- electronic mail
- DNS
- **P2P applications**
- video streaming and content distribution networks


79

Pure P2P architecture

- *no* always-on server
- arbitrary end systems directly communicate
- peers are intermittently connected and change IP addresses

examples:


- file distribution (BitTorrent)
- Streaming (KanKan)
- VoIP (Skype)



80

File distribution: client-server vs P2P

Question: how much time to distribute file (size F) from one server to N peers?
– peer upload/download capacity is limited resource



81

File distribution time: client-server

- **server transmission:** must sequentially send (upload) N file copies:

- time to send one copy: F/u_s
- time to send N copies: NF/u_s

- **client:** each client must download file copy

- d_{\min} = min client download rate
- min client download time: F/d_{\min}

$$\text{time to distribute } F \text{ to } N \text{ clients using client-server approach} \quad D_{c-s} \geq \max\{NF/u_s, F/d_{\min}\}$$

increases linearly in N



82

File distribution time: P2P

- **server transmission:** must upload at least one copy
- time to send one copy: F/u_s

- **client:** each client must download file copy


- min client download time: F/d_{\min}

- **clients:** as aggregate must download NF bits

- max upload rate (limiting max download rate) is $u_s + \sum u_i$

$$\text{time to distribute } F \text{ to } N \text{ clients using P2P approach} \quad D_{P2P} \geq \max\{F/u_s, F/d_{\min}, NF/(u_s + \sum u_i)\}$$


increases linearly in N ...
... but so does this, as each peer brings service capacity



83

Client-server vs. P2P: example

client upload rate = u , $F/u = 1$ hour, $u_s = 10u$, $d_{\min} \geq u_s$



84

P2P file distribution: BitTorrent

- file divided into 256Kb chunks
 - peers in torrent send/receive file chunks
 - tracker**: tracks peers participating in torrent
 - torrent**: group of peers exchanging chunks of a file
- Alice arrives ...
... obtains list of peers from tracker
... and begins exchanging file chunks with peers in torrent
-

85

P2P file distribution: BitTorrent

- peer joining torrent:
 - has no chunks, but will accumulate them over time from other peers
 - registers with tracker to get list of peers, connects to subset of peers ("neighbors")
 - while downloading, peer uploads chunks to other peers
 - peer may change peers with whom it exchanges chunks
 - churn**: peers may come and go
 - once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent
-

86


BitTorrent: requesting, sending file chunks

- requesting chunks:**
- at any given time, different peers have different subsets of file chunks
 - periodically, Alice asks each peer for list of chunks that they have
 - Alice requests missing chunks from peers, rarest first
- sending chunks: tit-for-tat**
- Alice sends chunks to those four peers currently sending her chunks *at highest rate*
 - other peers are choked by Alice (do not receive chunks from her)
 - re-evaluate top 4 every 10 secs
 - every 30 secs: randomly select another peer, starts sending chunks
 - "optimistically unchoke" this peer
 - newly chosen peer may join top 4

87

BitTorrent: tit-for-tat

- (1) Alice "optimistically unchokes" Bob
- (2) Alice becomes one of Bob's top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice's top-four providers



88

Video Streaming and CDNs: context

- video traffic: major consumer of Internet bandwidth
 - Netflix, YouTube: 37%, 16% of downstream residential ISP traffic
 - ~1B YouTube users, ~75M Netflix users
- challenge: scale - how to reach ~1B users?
 - single mega-video server won't work (why?)
- challenge: heterogeneity
 - different users have different capabilities (e.g., wired versus mobile; bandwidth rich versus bandwidth poor)
- solution**: distributed, application-level infrastructure



89

Multimedia: video


- video: sequence of images displayed at constant rate
 - e.g., 24 images/sec
- digital image: array of pixels
 - each pixel represented by bits
- coding: use redundancy *within* and *between* images to decrease # bits used to encode image
 - spatial (within image)
 - temporal (from one image to next)



90

Multimedia: video


- **CBR: (constant bit rate):** video encoding rate fixed
- **VBR: (variable bit rate):** video encoding rate changes as amount of spatial, temporal coding changes
- **examples:**
 - MPEG 1 (CD-ROM) 1.5 Mbps
 - MPEG2 (DVD) 3-6 Mbps
 - MPEG4 (often used in Internet, < 1 Mbps)



91

Streaming stored video:

simple scenario:



92

Streaming multimedia: DASH

- **DASH: Dynamic, Adaptive Streaming over HTTP**
- **server:**
 - divides video file into multiple chunks
 - each chunk stored, encoded at different rates
 - **manifest file:** provides URLs for different chunks
- **client:**
 - periodically measures server-to-client bandwidth
 - consulting manifest, requests one chunk at a time
 - chooses maximum coding rate sustainable given current bandwidth
 - can choose different coding rates at different points in time (depending on available bandwidth at time)

93

Streaming multimedia: DASH

- **DASH: Dynamic, Adaptive Streaming over HTTP**
- **"Intelligence" at client:** client determines
 - **when** to request chunk (so that buffer starvation, or overflow does not occur)
 - **what encoding rate** to request (higher quality when more bandwidth available)
 - **where** to request chunk (can request from URL server that is "close" to client or has high available bandwidth)

94

Content distribution networks

- **challenge:** how to stream content (selected from millions of videos) to hundreds of thousands of *simultaneous* users?
- **option 1:** single, large "mega-server"
 - single point of failure
 - point of network congestion
 - long path to distant clients
 - multiple copies of video sent over outgoing link

....quite simply: this solution *doesn't scale*

95


Content distribution networks

- **challenge:** how to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users?
- **option 2:** store/serve multiple copies of videos at multiple geographically distributed sites (**CDN**)
 - **enter deep:** push CDN servers deep into many access networks
 - close to users
 - used by Akamai, 1700 locations
 - **bring home:** smaller number (10's) of larger clusters in POPs near (but not within) access networks
 - used by Limelight

96

Content distribution networks


- CDN: stores copies of content at CDN nodes
 - e.g. Netflix stores copies of MadMen
- subscriber requests content from CDN
 - directed to nearby copy, retrieves content
 - may choose different copy if network path congested



97


CDN content access: a closer look

Bob (client) requests video <http://netcinema.com/6Y7B23V>
 video stored in CDN at <http://KingCDN.com/NetC6y&B23V>



98

Case study: Netflix



99